

▼ PART-C

▼ Q4

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import metrics
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import roc_curve, auc, confusion_matrix, classification_report
from sklearn.ensemble import RandomForestClassifier
from datascience import *
from sklearn.model_selection import cross_val_score
from sklearn import tree
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
%matplotlib inline
```

▼ Data Validation

```
data=pd.read_csv("/content/ckd_full.csv")
data.head()
```

| | Age | Blood Pressure | Specific Gravity | Albumin | Sugar | Red Blood Cells | Pus Cell | Pus Cell clumps | Bacteria | Blood Glucose Random | ... |
|---|------|----------------|------------------|---------|-------|-----------------|----------|-----------------|------------|----------------------|-----|
| 0 | 48.0 | 80.0 | 1.020 | 1.0 | 0.0 | NaN | normal | notpresent | notpresent | 121.0 | ... |
| 1 | 7.0 | 50.0 | 1.020 | 4.0 | 0.0 | NaN | normal | notpresent | notpresent | NaN | ... |
| 2 | 62.0 | 80.0 | 1.010 | 2.0 | 3.0 | normal | normal | notpresent | notpresent | 423.0 | ... |
| 3 | 48.0 | 70.0 | 1.005 | 4.0 | 0.0 | normal | abnormal | present | notpresent | 117.0 | ... |
| 4 | 51.0 | 80.0 | 1.010 | 2.0 | 0.0 | normal | normal | notpresent | notpresent | 106.0 | ... |

5 rows × 25 columns



#Data Cleaning

```
data['Hypertension'] = data['Hypertension'].replace(to_replace={'yes':1,'no':0})
data['Diabetes Mellitus'] = data['Diabetes Mellitus'].replace(to_replace={'yes':1,'no':0})
data['Coronary Artery Disease'] = data['Coronary Artery Disease'].replace(to_replace={'yes':1,'no':0})
data['Appetite'] = data['Appetite'].replace(to_replace={'good':1,'poor':0})
data['Pedal Edema'] = data['Pedal Edema'].replace(to_replace={'yes':1,'no':0})
data['Anemia'] = data['Anemia'].replace(to_replace={'yes':1,'no':0})
```

```
data['Red Blood Cells'] = data['Red Blood Cells'].replace(to_replace={'abnormal':1,'normal':0})
data['Pus Cell'] = data['Pus Cell'].replace(to_replace={'abnormal':1,'normal':0})
```

```
data['Pus Cell clumps'] = data['Pus Cell clumps'].replace(to_replace={'present':1,'notpresent':0})
data['Bacteria'] = data['Bacteria'].replace(to_replace={'present':1,'notpresent':0})
```

```
data['Class'] = data['Class'].replace(to_replace={'ckd':1.0,'ckd\t':1.0,'notckd':0.0})
```

```
data.to_csv("Out1.csv")
```

```
data['Pedal Edema'] = data['Pedal Edema'].replace(to_replace='good',value=0)
data['Appetite'] = data['Appetite'].replace(to_replace='no',value=0)
data['Coronary Artery Disease'] = data['Coronary Artery Disease'].replace(to_replace='yes',value=1)
data['Diabetes Mellitus'] = data['Diabetes Mellitus'].replace(to_replace={'\t':0,'yes':1,'no':0})
```

```
data1=data.dropna(axis=0)
```

```
data1.shape
```

```
(158, 25)
```

```
data1.isna().sum().sort_values(ascending=False)
```

| | |
|-------------------------|---|
| Age | 0 |
| Potassium | 0 |
| Anemia | 0 |
| Pedal Edema | 0 |
| Appetite | 0 |
| Coronary Artery Disease | 0 |
| Diabetes Mellitus | 0 |
| Hypertension | 0 |
| Red Blood Cell Count | 0 |
| White Blood Cell Count | 0 |
| Packed Cell Volume | 0 |
| Hemoglobin | 0 |
| Sodium | 0 |
| Blood Pressure | 0 |
| Serum Creatinine | 0 |
| Blood Urea | 0 |
| Blood Glucose Random | 0 |
| Bacteria | 0 |
| Pus Cell clumps | 0 |

```

Pus Cell      0
Red Blood Cells 0
Sugar         0
Albumin       0
Specific Gravity 0
Class        0
dtype: int64

```

```

corr_df = data1.corr()
mask = np.zeros_like(corr_df, dtype=np.bool)
mask[np.triu_indices_from(mask)] = True
f, ax = plt.subplots(figsize=(11, 9))
cmap = sns.diverging_palette(220, 10, as_cmap=True)
sns.heatmap(corr_df, mask=mask, cmap=cmap, vmax=.3, center=0,
            square=True, linewidths=.5, cbar_kws={"shrink": .5})
plt.title('Correlations between different predictors')
plt.show()


```

<ipython-input-7-27a1ceaba938>:2: DeprecationWarning: `np.bool` is a deprecated alias for the Deprecated in NumPy 1.20; for more details and guidance: <https://numpy.org/devdocs/release/1.2>

```

mask = np.zeros_like(corr_df, dtype=np.bool)

```



▼ KNN

```
def standard_units(x):
    return (x - np.mean(x))/np.std(x)

ckd = pd.DataFrame(
    {'Hemoglobin':standard_units(data1['Hemoglobin']),
    'Glucose':standard_units(data1['Blood Glucose Random']),
    'Class':data1['Class']}
)

color_table = pd.DataFrame(
    {'Class':np.array([1, 0]),
    'Color':np.array(['darkblue', 'gold'])}
)

ckd = pd.merge(ckd, color_table, on='Class')

ckd
```

| | Hemoglobin | Glucose | Class | Color |
|-----|------------|-----------|-------|----------|
| 0 | -0.865744 | -0.221549 | 1.0 | darkblue |
| 1 | -1.457446 | -0.947597 | 1.0 | darkblue |
| 2 | -1.004968 | 3.841231 | 1.0 | darkblue |
| 3 | -2.814879 | 0.396364 | 1.0 | darkblue |
| 4 | -2.083954 | 0.643529 | 1.0 | darkblue |
| ... | ... | ... | ... | ... |
| 153 | 0.700526 | 0.133751 | 0.0 | gold |
| 154 | 0.978974 | -0.870358 | 0.0 | gold |
| 155 | 0.735332 | -0.484162 | 0.0 | gold |
| 156 | 0.178436 | -0.267893 | 0.0 | gold |
| 157 | 0.735332 | -0.005280 | 0.0 | gold |

158 rows × 4 columns

```
#Alice in Scatter plot
alice = np.array([0, 1.1])

ckd_darkblue = ckd[ckd['Color'] == 'darkblue']
ckd_gold = ckd[ckd['Color'] == 'gold']


fig, ax = plt.subplots(figsize=(6,6))
```

```

ax.scatter(ckd_darkblue['Hemoglobin'],
           ckd_darkblue['Glucose'],
           label='Color=darkblue',
           color='darkblue')
ax.scatter(ckd_gold['Hemoglobin'],
           ckd_gold['Glucose'],
           label='Color=gold',
           color='gold')
ax.scatter(alice[0],
           alice[1],
           color='red',
           s=30)

y_vals = ax.get_yticks()
plt.ylabel('Glucose')
ax.legend(bbox_to_anchor=(1.04,1), loc="upper left")
plt.xlabel('Hemoglobin')
plt.xlim(-4, 2)
plt.ylim(-2, 6);
plt.show()

```



```

#Euclidean Distance
def distance(point1, point2):
    return np.sqrt(np.sum((point1 - point2)**2))

def distance_from_alice(row):
    return distance(alice, np.array(row))

```

```
ckd1=ckd[['Hemoglobin','Glucose']]
ckd1
```

| | Hemoglobin | Glucose |
|-----|------------|-----------|
| 0 | -0.865744 | -0.221549 |
| 1 | -1.457446 | -0.947597 |
| 2 | -1.004968 | 3.841231 |
| 3 | -2.814879 | 0.396364 |
| 4 | -2.083954 | 0.643529 |
| ... | ... | ... |
| 153 | 0.700526 | 0.133751 |
| 154 | 0.978974 | -0.870358 |
| 155 | 0.735332 | -0.484162 |
| 156 | 0.178436 | -0.267893 |
| 157 | 0.735332 | -0.005280 |

158 rows × 2 columns

```
distances = ckd1.apply(distance_from_alice, axis=1)
ckd_with_distances = ckd.copy()
ckd_with_distances['Distance from Alice'] = distances
```

ckd_with_distances

| | Hemoglobin | Glucose | Class | Color | Distance from Alice |
|-----|------------|-----------|-------|----------|---------------------|
| 0 | -0.865744 | -0.221549 | 1.0 | darkblue | 1.579875 |
| 1 | -1.457446 | -0.947597 | 1.0 | darkblue | 2.513325 |
| 2 | -1.004968 | 3.841231 | 1.0 | darkblue | 2.919641 |
| 3 | -2.814879 | 0.396364 | 1.0 | darkblue | 2.901491 |
| 4 | -2.083954 | 0.643529 | 1.0 | darkblue | 2.133361 |
| ... | ... | ... | ... | ... | ... |
| 153 | 0.700526 | 0.133751 | 0.0 | gold | 1.193471 |
| 154 | 0.978974 | -0.870358 | 0.0 | gold | 2.200159 |
| 155 | 0.735332 | -0.484162 | 0.0 | gold | 1.746506 |
| 156 | 0.178436 | -0.267893 | 0.0 | gold | 1.379482 |
| 157 | 0.735332 | -0.005280 | 0.0 | gold | 1.327537 |

158 rows × 5 columns

```
sorted_by_distance = ckd_with_distances.sort_values(by=['Distance from Alice'])
sorted_by_distance
```

| | Hemoglobin | Glucose | Class | Color | Distance from Alice |
|-----|------------|-----------|-------|----------|---------------------|
| 14 | 0.839750 | 1.215099 | 1.0 | darkblue | 0.847601 |
| 35 | -0.970162 | 1.276890 | 1.0 | darkblue | 0.986156 |
| 84 | -0.030400 | 0.087407 | 0.0 | gold | 1.013049 |
| 152 | 0.143630 | 0.087407 | 0.0 | gold | 1.022728 |
| 6 | -0.413266 | 2.049282 | 1.0 | darkblue | 1.035338 |
| ... | ... | ... | ... | ... | ... |
| 2 | -1.004968 | 3.841231 | 1.0 | darkblue | 2.919641 |
| 12 | -2.292790 | -0.854910 | 1.0 | darkblue | 3.013065 |
| 41 | -0.378460 | 4.520935 | 1.0 | darkblue | 3.441806 |
| 42 | -3.685029 | 0.689873 | 1.0 | darkblue | 3.707782 |
| 36 | -0.761326 | 5.540492 | 1.0 | darkblue | 4.505285 |

158 rows × 5 columns

```
alice_5_nearest_neighbors = sorted_by_distance.take(np.arange(5))
alice_5_nearest_neighbors
```

| | Hemoglobin | Glucose | Class | Color | Distance from Alice |
|-----|------------|----------|-------|----------|---------------------|
| 14 | 0.839750 | 1.215099 | 1.0 | darkblue | 0.847601 |
| 35 | -0.970162 | 1.276890 | 1.0 | darkblue | 0.986156 |
| 84 | -0.030400 | 0.087407 | 0.0 | gold | 1.013049 |
| 152 | 0.143630 | 0.087407 | 0.0 | gold | 1.022728 |
| 6 | -0.413266 | 2.049282 | 1.0 | darkblue | 1.035338 |

```
a=alice_5_nearest_neighbors.groupby('Class').count()
```

```
nothaving_ckd=a.iloc[0,3]
nothaving_ckd
```

2

```
Having_ckd=a.iloc[1,3]
Having_ckd
```

3

```
if nothaving_ckd>Having_ckd:
    print("Alice does not have Chronic Kidney disease")
```

```
else:
    print("Alice has Chronic Kidney disease")

    Alice has Chronic Kidney disease
```

RandomForest

```
data1.head()
```

| | Age | Blood Pressure | Specific Gravity | Albumin | Sugar | Red Blood Cells | Pus Cell | Pus Cell clumps | Bacteria | Blood Glucose Random | ... | Packed Cell Volume |
|----|------|----------------|------------------|---------|-------|-----------------|----------|-----------------|----------|----------------------|-----|--------------------|
| 3 | 48.0 | 70.0 | 1.005 | 4.0 | 0.0 | 0.0 | 1.0 | 1.0 | 0.0 | 117.0 | ... | 32 |
| 9 | 53.0 | 90.0 | 1.020 | 2.0 | 0.0 | 1.0 | 1.0 | 1.0 | 0.0 | 70.0 | ... | 29 |
| 11 | 63.0 | 70.0 | 1.010 | 3.0 | 0.0 | 1.0 | 1.0 | 1.0 | 0.0 | 380.0 | ... | 32 |
| 14 | 68.0 | 80.0 | 1.010 | 3.0 | 2.0 | 0.0 | 1.0 | 1.0 | 1.0 | 157.0 | ... | 16 |
| 20 | 61.0 | 80.0 | 1.015 | 2.0 | 0.0 | 1.0 | 1.0 | 0.0 | 0.0 | 173.0 | ... | 24 |

5 rows × 25 columns



```
ckd.head()
```

| | Hemoglobin | Glucose | Class | Color |
|---|------------|-----------|-------|----------|
| 0 | -0.865744 | -0.221549 | 1.0 | darkblue |
| 1 | -1.457446 | -0.947597 | 1.0 | darkblue |
| 2 | -1.004968 | 3.841231 | 1.0 | darkblue |
| 3 | -2.814879 | 0.396364 | 1.0 | darkblue |
| 4 | -2.083954 | 0.643529 | 1.0 | darkblue |



```
ckd_copy=ckd.copy()
ckd_copy=ckd_copy.drop('Color',1)
```

```
<ipython-input-24-31e986d8ffe8>:2: FutureWarning: In a future version of pandas all arguments
ckd_copy=ckd_copy.drop('Color',1)
```



```
x=ckd_copy.drop('Class',axis=1)
y=ckd_copy['Class']
X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=42)
print("X_train size {} , X_test size {}".format(X_train.shape,X_test.shape))
```


X_train size (126, 2) , X_test size (32, 2)

```
score=cross_val_score(RandomForestClassifier(max_depth=15,n_estimators=5),X_train,y_train,scoring='accuracy')
print("Average Accuracy Score {}".format(score.mean()))
```

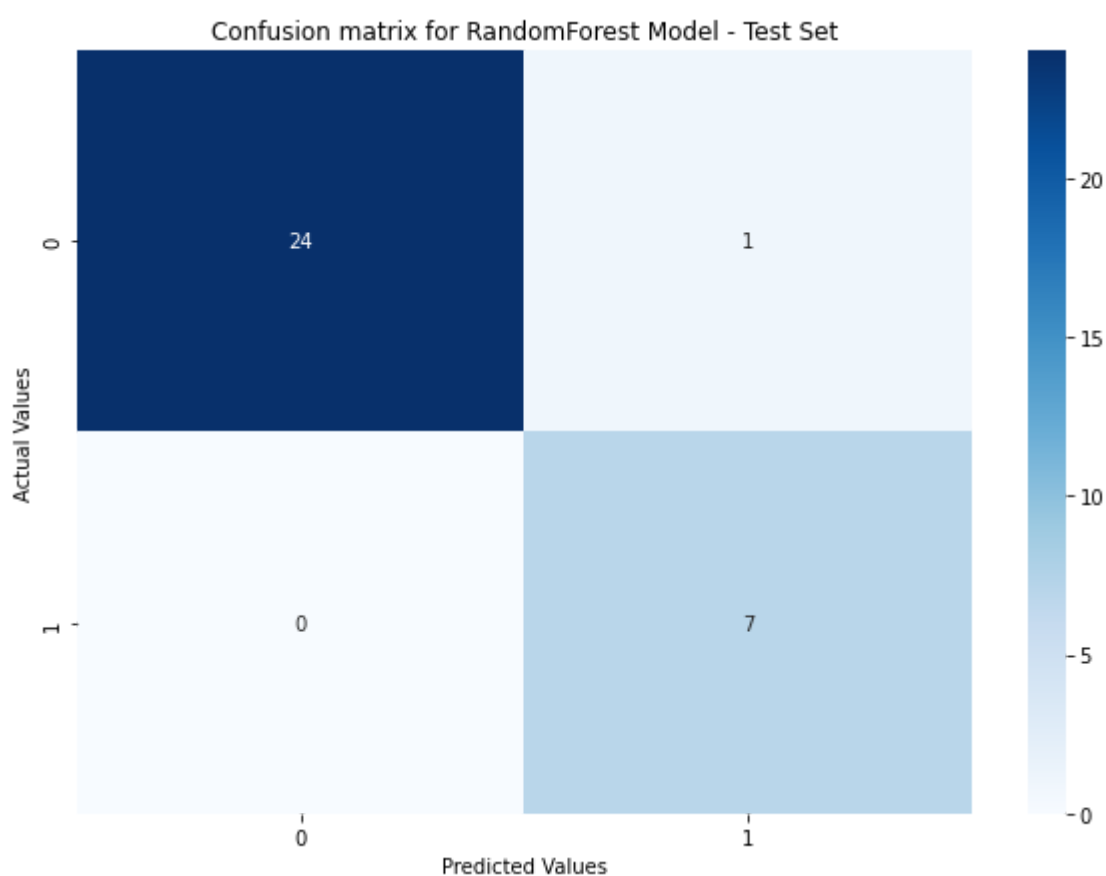
Average Accuracy Score 0.9923076923076923

```
rf=RandomForestClassifier(max_depth=5,n_estimators=5)
rf.fit(X_train,y_train)
```

RandomForestClassifier(max_depth=5, n_estimators=5)

```
y_pred=rf.predict(X_test)
confusionmatrix=confusion_matrix(y_pred,y_test)
```

```
plt.figure(figsize=(10,7))
p = sns.heatmap(confusionmatrix, annot=True, cmap="Blues", fmt='g')
plt.title('Confusion matrix for RandomForest Model - Test Set')
plt.xlabel('Predicted Values')
plt.ylabel('Actual Values')
plt.show()
```



```
score=round(accuracy_score(y_test,y_pred),3)
print("Accuracy on the Test set: {}".format(score))
```

Accuracy on the Test set: 0.969

```
print(classification_report(y_test,y_pred))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0 | 0.96 | 1.00 | 0.98 | 24 |
| 1.0 | 1.00 | 0.88 | 0.93 | 8 |
| accuracy | | | 0.97 | 32 |
| macro avg | 0.98 | 0.94 | 0.96 | 32 |
| weighted avg | 0.97 | 0.97 | 0.97 | 32 |

```
X_train=X_train[['Hemoglobin','Glucose']]
```

```
X_test=X_test[['Hemoglobin','Glucose']]
```

```
rf.fit(X_train,y_train)
```

```
def predict(hemo,gl):
```

```
    x=[[hemo,gl]]
```

```
    return rf.predict(x)
```

```
prediction=predict(0,1.1)[0]
```

```
prediction
```

```
/usr/local/lib/python3.8/dist-packages/sklearn/base.py:450: UserWarning: X does not have valid
warnings.warn(
1.0
```

```
if prediction:
```

```
    print('Oops! You have Chronic Kidney Disease.')
```

```
else:
```

```
    print("Great! You don't have Chronic Kidney Disease.")
```

```
    Oops! You have Chronic Kidney Disease.
```

DecisionTree

```
ckd_copy2=ckd_copy.copy()
```


```
ckd_copy2
```

| | Hemoglobin | Glucose | Class |  |
|---|------------|-----------|-------|---|
| 0 | -0.865744 | -0.221549 | 1.0 | |
| 1 | -1.457446 | -0.947597 | 1.0 | |
| 2 | -1.004968 | 3.841231 | 1.0 | |
| 3 | -2.814879 | 0.396364 | 1.0 | |

```
dtc = tree.DecisionTreeClassifier()
model = dtc.fit(X_train, y_train)
```

```
153    0.700526    0.133751    0.0
```

X_train

| | Hemoglobin | Glucose |  |
|-----|------------|-----------|---|
| 16 | -1.561864 | 2.528165 | |
| 130 | 0.874556 | -0.746776 | |
| 134 | 0.909362 | -0.314236 | |
| 22 | -2.014342 | 2.420030 | |
| 93 | 0.074018 | -0.345132 | |
| ... | ... | ... | |
| 9 | -1.353028 | 0.489051 | |
| 103 | 1.431451 | 0.010168 | |
| 67 | 0.282854 | -0.298788 | |
| 117 | 0.456884 | -0.561402 | |
| 47 | -0.239236 | -0.499610 | |

126 rows × 2 columns

```
predict_train = model.predict(X_train)
predict_train
```

```
array([[ 1.,  0.,  0.,  1.,  0.,  0.,  0.,  0.,  1.,  0.,  0.,  0.,  1.,
         0.,  1.,  0.,  0.,  1.,  0.,  0.,  0.,  1.,  0.,  0.,  0.,  0.,
         0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,
         1.,  0.,  0.,  1.,  0.,  0.,  0.,  0.,  1.,  0.,  1.,  0.,  0.,
         0.,  0.,  1.,  0.,  0.,  1.,  0.,  1.,  0.,  0.,  1.,  0.,  1.,
         0.,  1.,  0.,  1.,  1.,  0.,  0.,  0.,  1.,  1.,  1.,  0.,  0.,
         0.,  1.,  1.,  0.,  0.,  1.,  0.,  0.,  0.,  1.,  1.,  0.,  0.,
         0.,  0.,  0.,  0.,  0.,  1.,  0.,  0.,  0.,  1.,  0.,  0.,  0.,
         0.,  0.,  1.,  1.,  1.,  0.,  0.,  0.,  0.]])
```

```
predict_test = model.predict(X_test)
predict_test
```

```
array([[ 1.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  1.,  0.,  1.,
         0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  1.,  0.,  1.,  1.,  1.]])
```

```
0., 1., 0., 0., 0., 0.]])
```

```
dtc_acc = accuracy_score(y_test, dtc.predict(X_test))
print("\n")
print(f"Training Accuracy of Decision Tree Classifier is {accuracy_score(y_train, dtc.predict(X_train))}")
print(f"Test Accuracy of Decision Tree Classifier is {dtc_acc} \n")
print("\n")
print(f"Confusion Matrix :- \n{confusion_matrix(y_test, dtc.predict(X_test))}")
confusion = confusion_matrix(y_test, dtc.predict(X_test))
tn, fp, fn, tp = confusion.ravel()
print("\n")

print("TN -->",tn)
print("FP -->",fp)
print("FN -->",fn)
print("TP -->",tp)
print("\n")
print("\n")

print(f"Classification Report :- \n {classification_report(y_test, dtc.predict(X_test))}")
```

```
Training Accuracy of Decision Tree Classifier is 1.0
Test Accuracy of Decision Tree Classifier is 1.0
```

```
Confusion Matrix :-
[[24  0]
 [ 0  8]]
```

```
TN --> 24
FP --> 0
FN --> 0
TP --> 8
```

```
Classification Report :-
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0 | 1.00 | 1.00 | 1.00 | 24 |
| 1.0 | 1.00 | 1.00 | 1.00 | 8 |
| accuracy | | | 1.00 | 32 |
| macro avg | 1.00 | 1.00 | 1.00 | 32 |
| weighted avg | 1.00 | 1.00 | 1.00 | 32 |

```
ind_col=ckd_copy2[['Hemoglobin','Glucose']]
dep_col=ckd_copy2['Class']
```

```

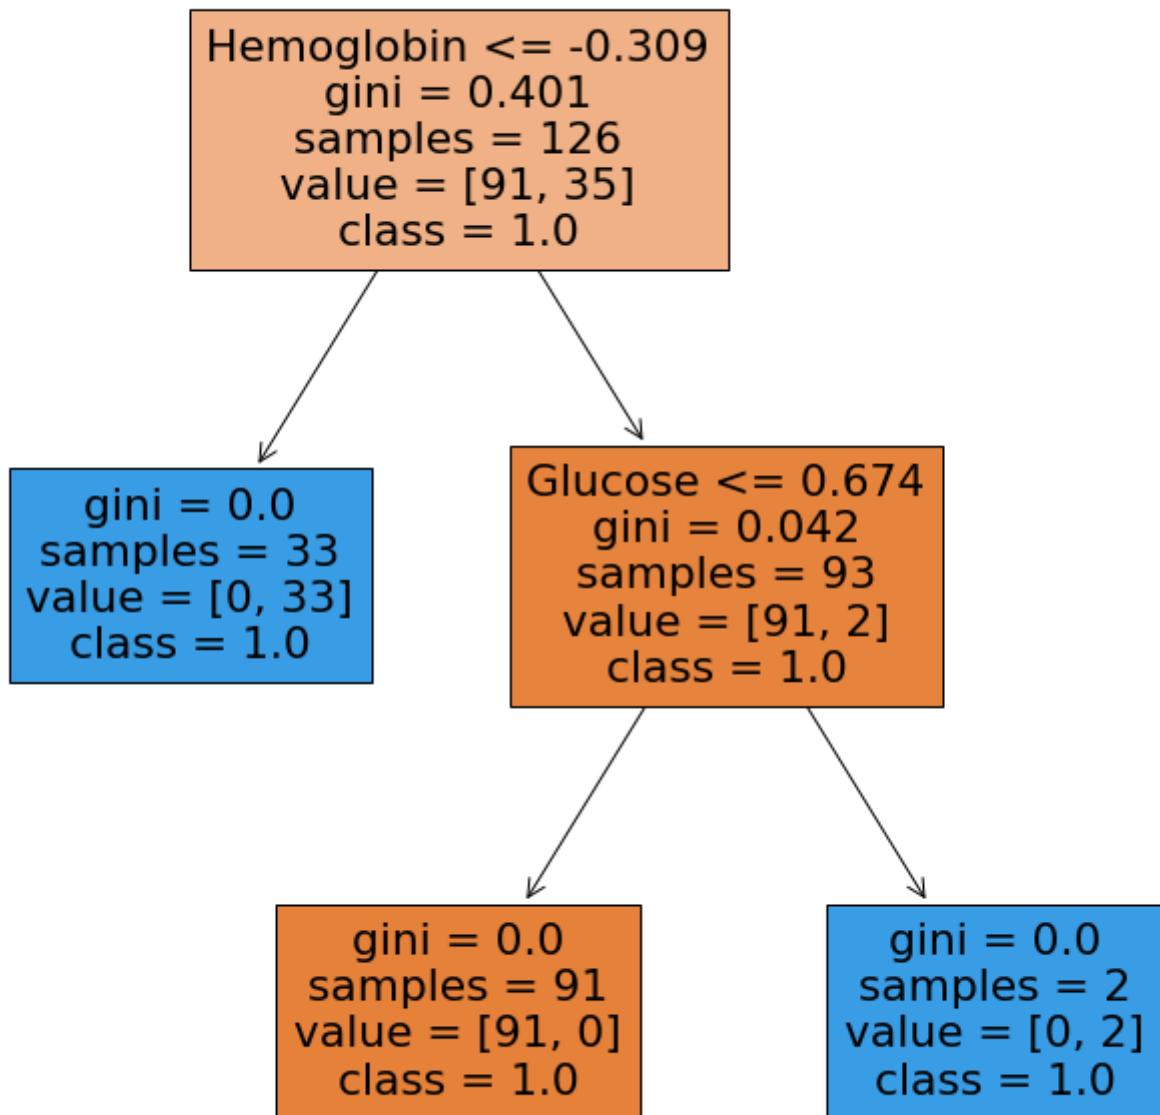
ind_col = ind_col.astype(str)
dep_col = dep_col.astype(str)
feature_names = ind_col.columns.tolist()
plt.figure(figsize=(12, 12))
plot_tree(model, filled=True, feature_names=feature_names, class_names=dep_col)

```

```

[Text(0.4, 0.83333333333333, 'Hemoglobin <= -0.309\ngini = 0.401\nsamples = 126\nvalue = [91, 35]\n= 1.0'),
Text(0.2, 0.5, 'gini = 0.0\nsamples = 33\nvalue = [0, 33]\nnclass = 1.0'),
Text(0.6, 0.5, 'Glucose <= 0.674\ngini = 0.042\nsamples = 93\nvalue = [91, 2]\nnclass = 1.0'),
Text(0.4, 0.16666666666667, 'gini = 0.0\nsamples = 91\nvalue = [91, 0]\nnclass = 1.0'),
Text(0.8, 0.16666666666667, 'gini = 0.0\nsamples = 2\nvalue = [0, 2]\nnclass = 1.0')]

```



```

Alice = np.array([[0, 1.1]])

```

```

final=model.predict(Alice)
final

```

```

/usr/local/lib/python3.8/dist-packages/sklearn/base.py:450: UserWarning: X does not have valid
warnings.warn(

```

```
array([ 1.])
```

```
if final==0:
    print("Alice does not have CKD")
else:
    print("Alice has CKD")
```

```
Alice has CKD
```

Q5.

Validation

```
wine=pd.read_csv("/content/winequality-red.csv")
wine.head()
```

| | fixed acidity | volatile acidity | citric acid | residual sugar | chlorides | free sulfur dioxide | total sulfur dioxide | density | pH | sulphates |
|---|---------------|------------------|-------------|----------------|-----------|---------------------|----------------------|---------|------|-----------|
| 0 | 7.4 | 0.70 | 0.00 | 1.9 | 0.076 | 11.0 | 34.0 | 0.9978 | 3.51 | 0.56 |
| 1 | 7.8 | 0.88 | 0.00 | 2.6 | 0.098 | 25.0 | 67.0 | 0.9968 | 3.20 | 0.68 |
| 2 | 7.8 | 0.76 | 0.04 | 2.3 | 0.092 | 15.0 | 54.0 | 0.9970 | 3.26 | 0.65 |
| 3 | 11.2 | 0.28 | 0.56 | 1.9 | 0.075 | 17.0 | 60.0 | 0.9980 | 3.16 | 0.58 |
| 4 | 7.4 | 0.70 | 0.00 | 1.9 | 0.076 | 11.0 | 34.0 | 0.9978 | 3.51 | 0.56 |



```
wine.isna().any()
```

```
fixed acidity      False
volatile acidity   False
citric acid        False
residual sugar     False
chlorides          False
free sulfur dioxide False
total sulfur dioxide False
density            False
pH                 False
sulphates          False
alcohol            False
quality            False
dtype: bool
```

```
wine['quality'].value_counts()
```

```
5    681
6    638
7    199
```

```

4      53
8      18
3      10
Name: quality, dtype: int64

```

```

corr_df = wine.corr()
mask = np.zeros_like(corr_df, dtype=np.bool)
mask[np.triu_indices_from(mask)] = True
f, ax = plt.subplots(figsize=(11, 9))
cmap = sns.diverging_palette(220, 10, as_cmap=True)
sns.heatmap(corr_df, mask=mask, cmap=cmap, vmax=.3, center=0,
            square=True, linewidths=.5, cbar_kws={"shrink": .5})
plt.title('Correlations between different predictors')
plt.show()

```

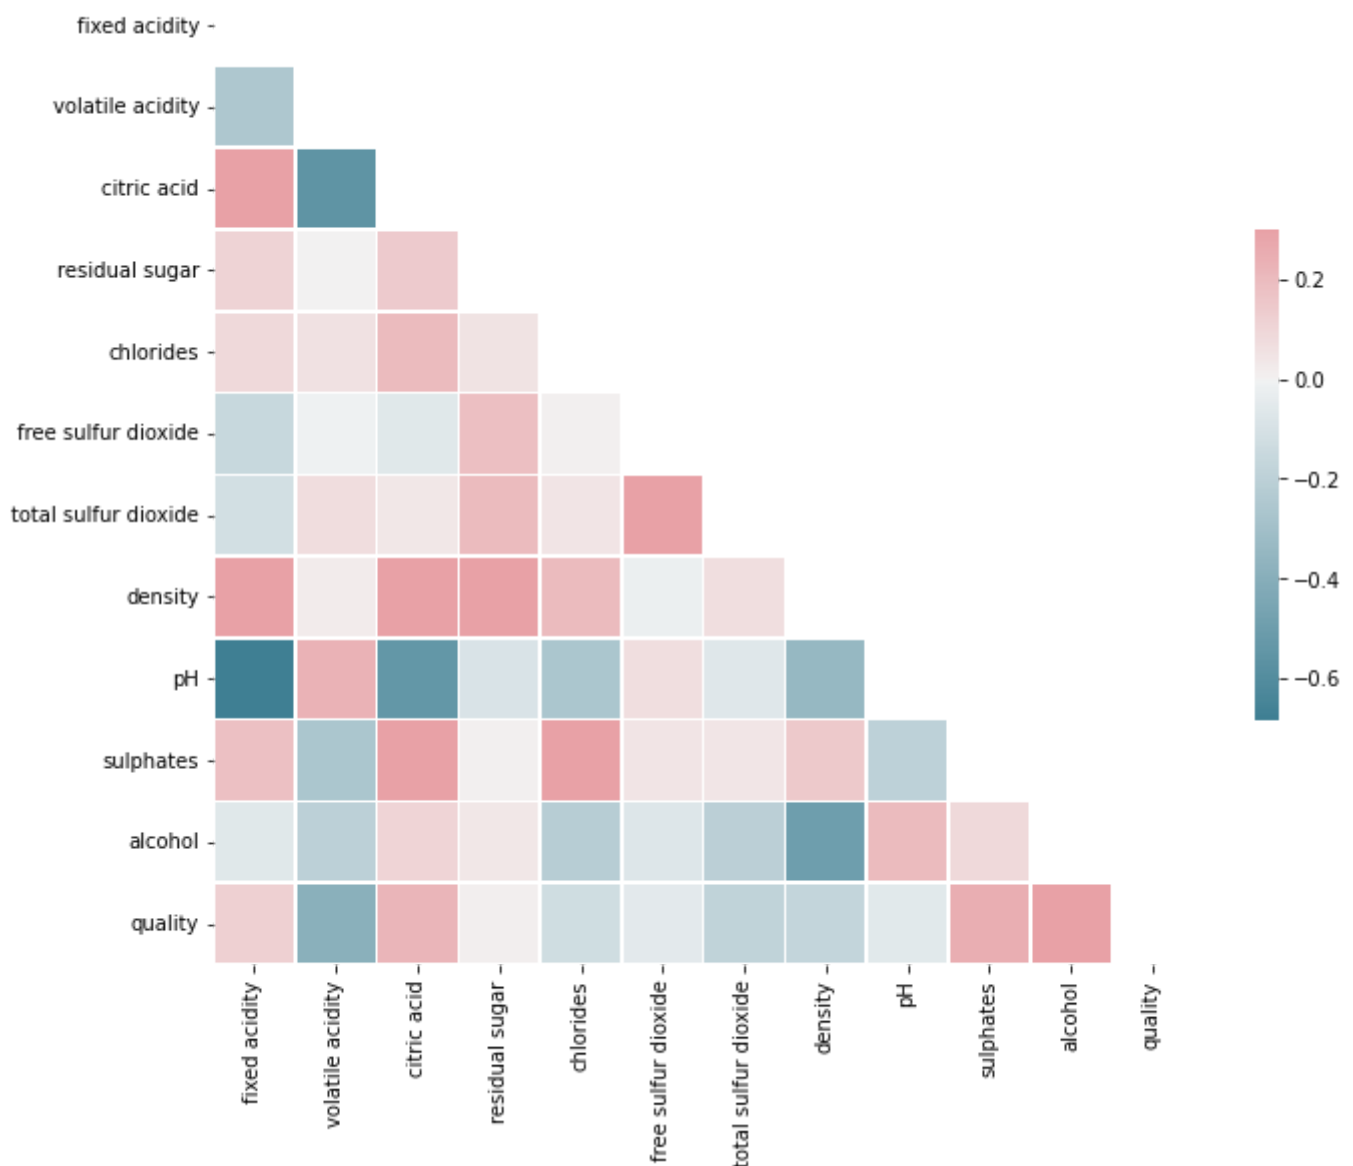
<ipython-input-47-af5de3b719fe>:2: DeprecationWarning: `np.bool` is a deprecated alias for the Deprecated in NumPy 1.20; for more details and guidance: <https://numpy.org/devdocs/release/1.20.0-notes>

```

mask = np.zeros_like(corr_df, dtype=np.bool)

```

Correlations between different predictors



```

wine.drop(["residual sugar", 'free sulfur dioxide', 'pH'], axis = 1, inplace = True)
wine.head()

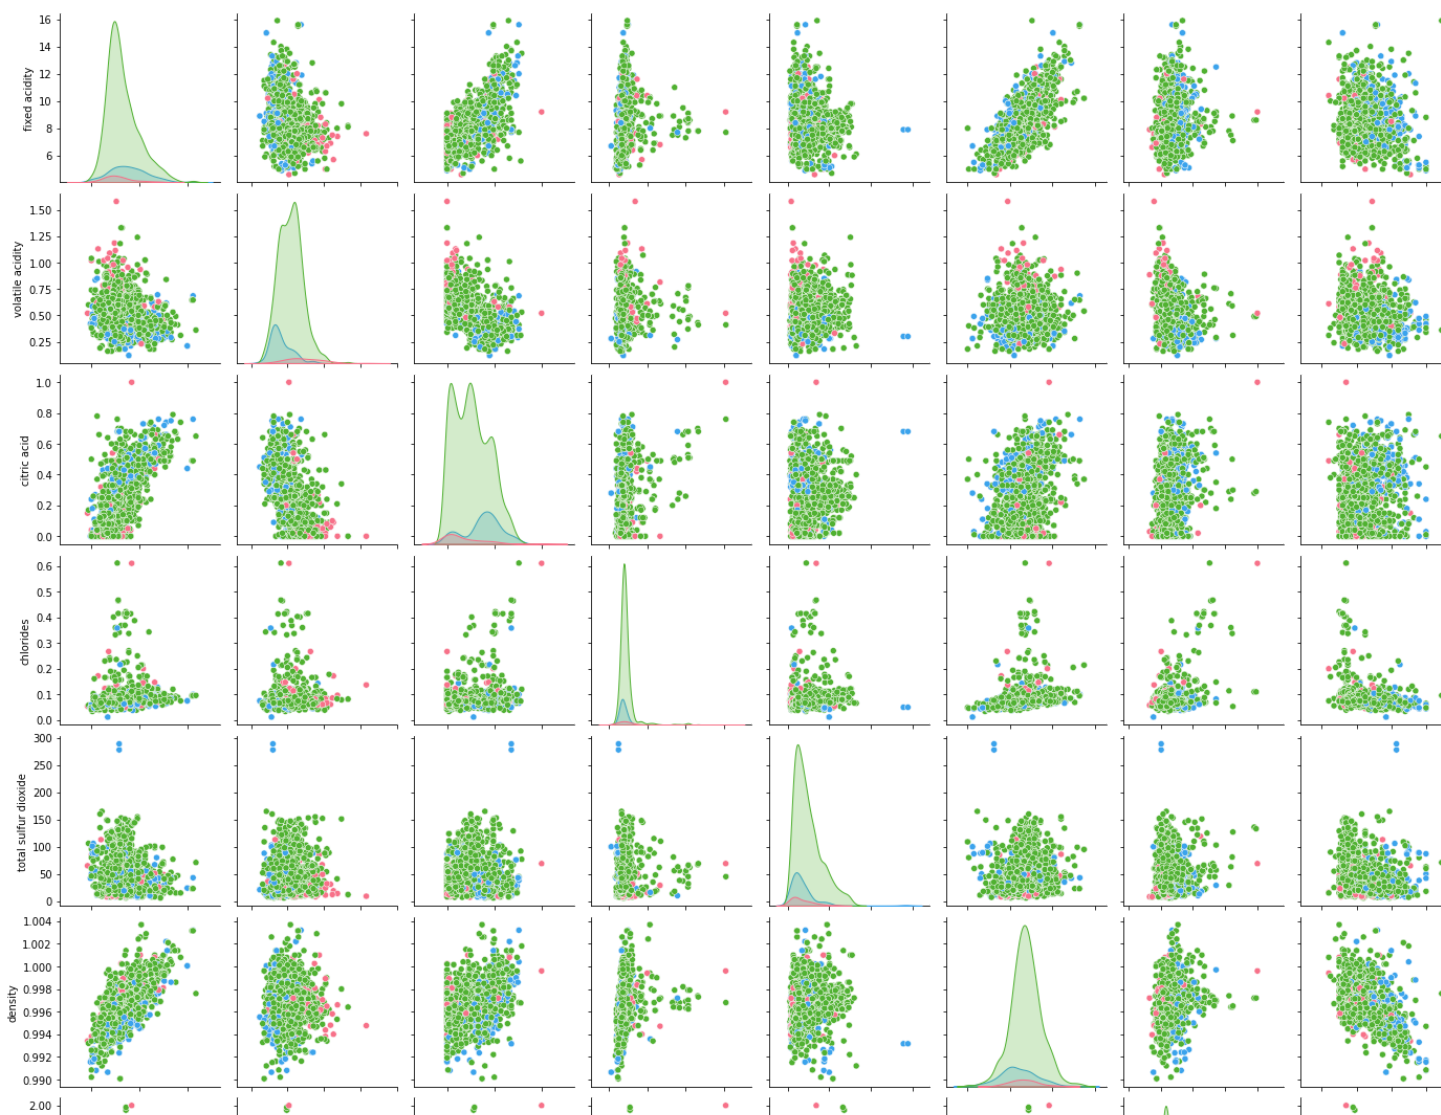
```

| | fixed acidity | volatile acidity | citric acid | chlorides | total sulfur dioxide | density | sulphates | alcohol |
|---|------------------|---------------------|----------------|-----------|-------------------------|---------|-----------|---------|
| 0 | 7.4 | 0.70 | 0.00 | 0.076 | 34.0 | 0.9978 | 0.56 | 9.4 |
| 1 | 7.8 | 0.88 | 0.00 | 0.098 | 67.0 | 0.9968 | 0.68 | 9.8 |
| 2 | 7.8 | 0.76 | 0.04 | 0.092 | 54.0 | 0.9970 | 0.65 | 9.8 |
| 3 | 11.2 | 0.28 | 0.56 | 0.075 | 60.0 | 0.9980 | 0.58 | 9.8 |
| 4 | 7.4 | 0.70 | 0.00 | 0.076 | 34.0 | 0.9978 | 0.56 | 9.4 |

```
bins = [0, 4, 6, 10]
labels = ["poor","normal","excellent"]
wine['quality_label'] = pd.cut(wine['quality'], bins=bins, labels=labels)
wine.drop('quality',axis =1, inplace = True)
wine.head()
```

| | fixed acidity | volatile acidity | citric acid | chlorides | total sulfur dioxide | density | sulphates | alcohol | quality_label |
|---|------------------|---------------------|----------------|-----------|-------------------------|---------|-----------|---------|---------------|
| 0 | 7.4 | 0.70 | 0.00 | 0.076 | 34.0 | 0.9978 | 0.56 | 9.4 | poor |
| 1 | 7.8 | 0.88 | 0.00 | 0.098 | 67.0 | 0.9968 | 0.68 | 9.8 | poor |
| 2 | 7.8 | 0.76 | 0.04 | 0.092 | 54.0 | 0.9970 | 0.65 | 9.8 | poor |
| 3 | 11.2 | 0.28 | 0.56 | 0.075 | 60.0 | 0.9980 | 0.58 | 9.8 | poor |

```
sns.pairplot(wine, hue="quality_label", palette="husl",diag_kind="kde")
plt.show()
```

```
#df_wine = pd.get_dummies(wine, columns=['alcohol_label'], drop_first=True)
#df_wine.head()
```



```
wine.head()
```

| | fixed acidity | volatile acidity | citric acid | chlorides | total sulfur dioxide | density | sulphates | alcohol | quality |
|---|---------------|------------------|-------------|-----------|----------------------|---------|-----------|---------|---------|
| 0 | 7.4 | 0.70 | 0.00 | 0.076 | 34.0 | 0.9978 | 0.56 | 9.4 | |
| 1 | 7.8 | 0.88 | 0.00 | 0.098 | 67.0 | 0.9968 | 0.68 | 9.8 | |
| 2 | 7.8 | 0.76 | 0.04 | 0.092 | 54.0 | 0.9970 | 0.65 | 9.8 | |
| 3 | 11.2 | 0.28 | 0.56 | 0.075 | 60.0 | 0.9980 | 0.58 | 9.8 | |

```
wine_label=wine['quality_label']
#wine_rest=wine.drop(['alcohol_label','quality_label'],1)
wine_rest=wine.drop(['quality_label'],1)
```

<ipython-input-53-6eaaa66e9f66>:3: FutureWarning: In a future version of pandas all arguments wine_rest=wine.drop(['quality_label'],1)



```
wine_label
```

```
0      normal
1      normal
2      normal
3      normal
4      normal
...
1594   normal
1595   normal
1596   normal
1597   normal
1598   normal
Name: quality_label, Length: 1599, dtype: category
Categories (3, object): ['poor' < 'normal' < 'excellent']
```

wine_rest

| | fixed acidity | volatile acidity | citric acid | chlorides | total sulfur dioxide | density | sulphates |
|------|------------------|---------------------|----------------|-----------|-------------------------|---------|-----------|
| 0 | 7.4 | 0.700 | 0.00 | 0.076 | 34.0 | 0.99780 | 0.56 |
| 1 | 7.8 | 0.880 | 0.00 | 0.098 | 67.0 | 0.99680 | 0.68 |
| 2 | 7.8 | 0.760 | 0.04 | 0.092 | 54.0 | 0.99700 | 0.65 |
| 3 | 11.2 | 0.280 | 0.56 | 0.075 | 60.0 | 0.99800 | 0.58 |
| 4 | 7.4 | 0.700 | 0.00 | 0.076 | 34.0 | 0.99780 | 0.56 |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 1594 | 6.2 | 0.600 | 0.08 | 0.090 | 44.0 | 0.99490 | 0.58 |
| 1595 | 5.9 | 0.550 | 0.10 | 0.062 | 51.0 | 0.99512 | 0.76 |
| 1596 | 6.3 | 0.510 | 0.13 | 0.076 | 40.0 | 0.99574 | 0.75 |
| 1597 | 5.9 | 0.645 | 0.12 | 0.075 | 44.0 | 0.99547 | 0.71 |
| 1598 | 6.0 | 0.310 | 0.47 | 0.067 | 42.0 | 0.99549 | 0.66 |

1599 rows × 8 columns

▼ KNN

```
X_train, X_test, Y_train, Y_test = train_test_split(wine_rest, wine_label, test_size=0.2)
```

X_train

| | fixed acidity | volatile acidity | citric acid | chlorides | total sulfur dioxide | density | sulphates |
|------------|------------------|---------------------|----------------|-----------|-------------------------|---------|-----------|
| 730 | 9.5 | 0.550 | 0.66 | 0.387 | 37.0 | 0.99820 | 0.67 |
| 932 | 7.6 | 0.400 | 0.29 | 0.078 | 66.0 | 0.99710 | 0.59 |
| 821 | 4.9 | 0.420 | 0.00 | 0.048 | 42.0 | 0.99154 | 0.74 |
| 985 | 7.4 | 0.580 | 0.00 | 0.064 | 11.0 | 0.99562 | 0.58 |
| 549 | 9.0 | 0.530 | 0.49 | 0.171 | 25.0 | 0.99750 | 0.61 |
| ... | ... | ... | ... | ... | ... | ... | ... |

X_test

| | fixed acidity | volatile acidity | citric acid | chlorides | total sulfur dioxide | density | sulphates |
|-------------|------------------|---------------------|----------------|-----------|-------------------------|---------|-----------|
| 1429 | 7.9 | 0.180 | 0.40 | 0.049 | 67.0 | 0.99600 | 0.93 |
| 260 | 7.9 | 0.330 | 0.23 | 0.077 | 45.0 | 0.99625 | 0.65 |
| 916 | 5.3 | 0.715 | 0.19 | 0.161 | 62.0 | 0.99395 | 0.61 |
| 1141 | 8.2 | 0.380 | 0.32 | 0.080 | 71.0 | 0.99624 | 0.85 |
| 1574 | 5.6 | 0.310 | 0.78 | 0.074 | 92.0 | 0.99677 | 0.48 |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 298 | 7.2 | 0.650 | 0.02 | 0.094 | 31.0 | 0.99930 | 0.80 |
| 571 | 6.2 | 0.360 | 0.24 | 0.095 | 42.0 | 0.99460 | 0.57 |
| 605 | 8.3 | 0.600 | 0.13 | 0.085 | 24.0 | 0.99840 | 0.59 |
| 1548 | 11.2 | 0.400 | 0.50 | 0.099 | 50.0 | 0.99783 | 0.58 |
| 455 | 11.3 | 0.620 | 0.67 | 0.086 | 19.0 | 0.99880 | 0.69 |

480 rows × 8 columns

```

scaler = StandardScaler()
scaler.fit(wine_rest)
scaled_features = scaler.transform(wine_rest)
wine_rest_scaled= pd.DataFrame(scaled_features, columns=wine_rest.columns)

X_train_sc, X_test_sc, y_train_sc, y_test_sc = train_test_split(wine_rest_scaled, wine_rest_scaled['quality'],
                                                                    test_size=0.3, random_state=42)

X_train_sc = X_train_sc.to_numpy()
y_train_sc = y_train_sc.to_numpy()

def apply_knn(neigh, weight='uniform'):
    knn = KNeighborsClassifier(n_neighbors=neigh, weights=weight)
    knn.fit(X_train_sc,y_train_sc)
    pred_knn = knn.predict(X_test_sc)
    return pred_knn

```

```

model = KNeighborsClassifier()

params = {'n_neighbors':list(range(1, 50, 2)), 'weights':['uniform', 'distance']}

gs = GridSearchCV(model, params, cv = 5, n_jobs=-1)

gs_results = gs.fit(X_train_sc, y_train_sc)

print('Best Accuracy: ', gs_results.best_score_)
print('Best Parameters: ', gs_results.best_params_)

```

```

Best Accuracy: 0.852570467649
Best Parameters: {'n_neighbors': 9, 'weights': 'distance'}

```

```

pred_knn = apply_knn(9)
print('Accuracy of model at K=9 is', accuracy_score(y_test_sc, pred_knn))

```

```

Accuracy of model at K=9 is 0.84375
/usr/local/lib/python3.8/dist-packages/sklearn/base.py:443: UserWarning: X has feature names,
warnings.warn(

```



▼ RandomForest

```

rfc = RandomForestClassifier(n_estimators=200)
rfc.fit(X_train_sc, y_train_sc)
pred_rfc = rfc.predict(X_test_sc)

```

```

/usr/local/lib/python3.8/dist-packages/sklearn/base.py:443: UserWarning: X has feature names,
warnings.warn(

```



```
y_test_sc.shape
```

```
(480,)
```

```
pred_rfc.shape
```

```
(480,)
```

```
print(classification_report(y_test_sc, pred_rfc))
```

| | precision | recall | f1-score | support |
|-----------|-----------|--------|----------|---------|
| excellent | 0.82 | 0.55 | 0.66 | 65 |
| normal | 0.89 | 0.98 | 0.93 | 395 |
| poor | 0.00 | 0.00 | 0.00 | 20 |

| | | | |
|--------------|------|------|------|
| accuracy | | 0.88 | 480 |
| macro avg | 0.57 | 0.51 | 0.53 |
| weighted avg | 0.84 | 0.88 | 0.85 |

```
print(confusion_matrix(y_test_sc, pred_rfc))
```

```
[[ 36  29   0]
 [   8 386   1]
 [   0  20   0]]
```

the Accuracy of Random Forest is around 84%

Support Vector Classifier

```
svc = SVC()
svc.fit(X_train_sc, y_train_sc)
pred_svc = svc.predict(X_test_sc)
```

```
/usr/local/lib/python3.8/dist-packages/sklearn/base.py:443: UserWarning: X has feature names,
warnings.warn(
```

```
print(classification_report(y_test_sc, pred_svc))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| excellent | 0.89 | 0.25 | 0.39 | 65 |
| normal | 0.85 | 0.99 | 0.92 | 395 |
| poor | 0.00 | 0.00 | 0.00 | 20 |
| accuracy | | | 0.85 | 480 |
| macro avg | 0.58 | 0.41 | 0.43 | 480 |
| weighted avg | 0.82 | 0.85 | 0.81 | 480 |

```
/usr/local/lib/python3.8/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetri
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.8/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetri
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.8/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetri
_warn_prf(average, modifier, msg_start, len(result))
```

#Support vector classifier gets 79%

PART-D

Q6

```
house=pd.read_csv("/content/House.csv")
```

```
house.head()
```

| | price | area | bedrooms | bathrooms | stories | mainroad | guestroom | basement | hotwaterheating |
|---|----------|------|----------|-----------|---------|----------|-----------|----------|-----------------|
| 0 | 13300000 | 7420 | 4 | 2 | 3 | yes | no | no | no |
| 1 | 12250000 | 8960 | 4 | 4 | 4 | yes | no | no | no |
| 2 | 12250000 | 9960 | 3 | 2 | 2 | yes | no | yes | no |
| 3 | 12215000 | 7500 | 4 | 2 | 2 | yes | no | yes | no |
| 4 | 11410000 | 7420 | 4 | 1 | 2 | yes | yes | yes | no |



```
house.shape
```

(545, 13)

```
house.describe()
```

| | price | area | bedrooms | bathrooms | stories | parking |
|-------|--------------|--------------|------------|------------|------------|------------|
| count | 5.450000e+02 | 545.000000 | 545.000000 | 545.000000 | 545.000000 | 545.000000 |
| mean | 4.766729e+06 | 5150.541284 | 2.965138 | 1.286239 | 1.805505 | 0.693578 |
| std | 1.870440e+06 | 2170.141023 | 0.738064 | 0.502470 | 0.867492 | 0.861586 |
| min | 1.750000e+06 | 1650.000000 | 1.000000 | 1.000000 | 1.000000 | 0.000000 |
| 25% | 3.430000e+06 | 3600.000000 | 2.000000 | 1.000000 | 1.000000 | 0.000000 |
| 50% | 4.340000e+06 | 4600.000000 | 3.000000 | 1.000000 | 2.000000 | 0.000000 |
| 75% | 5.740000e+06 | 6360.000000 | 3.000000 | 2.000000 | 2.000000 | 1.000000 |
| max | 1.330000e+07 | 16200.000000 | 6.000000 | 4.000000 | 4.000000 | 3.000000 |



```
house.isna().any()
```

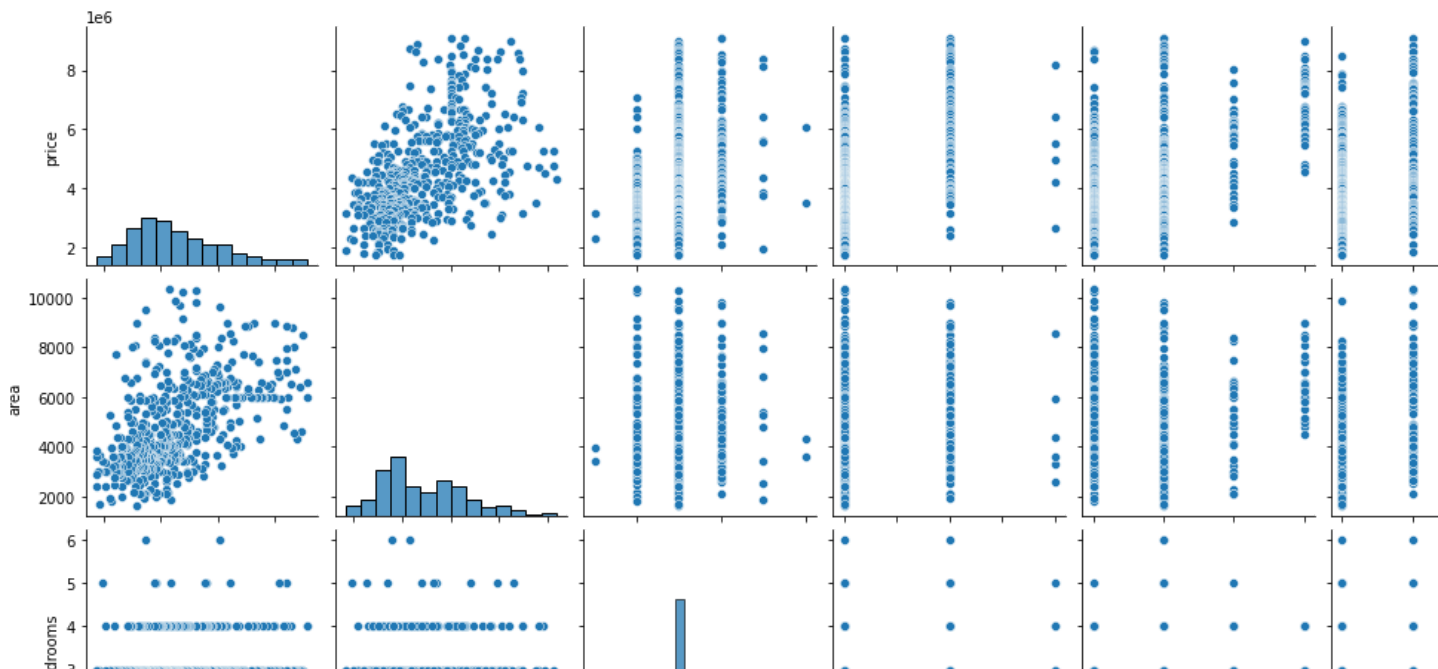
| | |
|-----------------|-------|
| price | False |
| area | False |
| bedrooms | False |
| bathrooms | False |
| stories | False |
| mainroad | False |
| guestroom | False |
| basement | False |
| hotwaterheating | False |
| airconditioning | False |
| parking | False |
| prefarea | False |

```
furnishingstatus      False  
dtype: bool
```

```
Q1 = house.price.quantile(0.25)  
Q3 = house.price.quantile(0.75)  
IQR = Q3 - Q1  
house = house[(house.price >= Q1 - 1.5*IQR) & (house.price <= Q3 + 1.5*IQR)]
```

```
Q1 = house.area.quantile(0.25)  
Q3 = house.area.quantile(0.75)  
IQR = Q3 - Q1  
house = house[(house.area >= Q1 - 1.5*IQR) & (house.area <= Q3 + 1.5*IQR)]
```

```
sns.pairplot(house)  
plt.show()
```



```
varlist = ['mainroad', 'guestroom', 'basement', 'hotwaterheating', 'airconditioning']
```

```
def binary_map(x):
    return x.map({'yes': 1, "no": 0})
```

```
house[varlist] = house[varlist].apply(binary_map)
```

```
house.head()
```

| | price | area | bedrooms | bathrooms | stories | mainroad | guestroom | basement | hotwaterheating |
|----|---------|------|----------|-----------|---------|----------|-----------|----------|-----------------|
| 15 | 9100000 | 6000 | 4 | 1 | 2 | 1 | 0 | 1 | |
| 16 | 9100000 | 6600 | 4 | 2 | 2 | 1 | 1 | 1 | |
| 17 | 8960000 | 8500 | 3 | 2 | 4 | 1 | 0 | 0 | |
| 18 | 8890000 | 4600 | 3 | 2 | 2 | 1 | 1 | 0 | |
| 19 | 8855000 | 6420 | 3 | 2 | 2 | 1 | 0 | 0 | |



```
status = pd.get_dummies(house['furnishingstatus'])
```

```
status.head()
```


furnished semi-furnished unfurnished 

```
status = pd.get_dummies(house['furnishingstatus'], drop_first = True)
house = pd.concat([house, status], axis = 1)
```

```
house.head()
```

| | price | area | bedrooms | bathrooms | stories | mainroad | guestroom | basement | hotwaterheating |
|----|---------|------|----------|-----------|---------|----------|-----------|----------|-----------------|
| 15 | 9100000 | 6000 | 4 | 1 | 2 | 1 | 0 | 1 | (|
| 16 | 9100000 | 6600 | 4 | 2 | 2 | 1 | 1 | 1 | (|
| 17 | 8960000 | 8500 | 3 | 2 | 4 | 1 | 0 | 0 | (|
| 18 | 8890000 | 4600 | 3 | 2 | 2 | 1 | 1 | 0 | (|
| 19 | 8855000 | 6420 | 3 | 2 | 2 | 1 | 0 | 0 | (|



```
house.drop(['furnishingstatus'], axis = 1, inplace = True)
house.head()
```

| | price | area | bedrooms | bathrooms | stories | mainroad | guestroom | basement | hotwaterheating |
|----|---------|------|----------|-----------|---------|----------|-----------|----------|-----------------|
| 15 | 9100000 | 6000 | 4 | 1 | 2 | 1 | 0 | 1 | (|
| 16 | 9100000 | 6600 | 4 | 2 | 2 | 1 | 1 | 1 | (|
| 17 | 8960000 | 8500 | 3 | 2 | 4 | 1 | 0 | 0 | (|
| 18 | 8890000 | 4600 | 3 | 2 | 2 | 1 | 1 | 0 | (|
| 19 | 8855000 | 6420 | 3 | 2 | 2 | 1 | 0 | 0 | (|



```
np.random.seed(0)
df_train, df_test = train_test_split(house, train_size = 0.7, test_size = 0.3)
```

```
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
num_vars = ['area', 'bedrooms', 'bathrooms', 'stories', 'parking', 'price']
df_train[num_vars] = scaler.fit_transform(df_train[num_vars])
df_train.head()
```

| | price | area | bedrooms | bathrooms | stories | mainroad | guestroom | basement | hotwaterh |
|------------|----------|----------|----------|-----------|----------|----------|-----------|----------|-----------|
| 148 | 0.523810 | 0.526907 | 0.4 | 0.0 | 0.666667 | 1 | 0 | 0 | |
| 236 | 0.390476 | 0.114134 | 0.2 | 0.0 | 0.333333 | 1 | 1 | 1 | |
| 356 | 0.275238 | 0.072738 | 0.8 | 0.5 | 0.000000 | 0 | 0 | 1 | |
| 425 | 0.219048 | 0.151390 | 0.2 | 0.0 | 0.000000 | 1 | 0 | 1 | |
| 516 | 0.095238 | 0.157895 | 0.2 | 0.0 | 0.000000 | 0 | 1 | 0 | |



```
df_train.describe()
```

| | price | area | bedrooms | bathrooms | stories | mainroad | guestroom | basement |
|--------------|------------|------------|------------|------------|------------|------------|------------|------------|
| count | 361.000000 | 361.000000 | 361.000000 | 361.000000 | 361.000000 | 361.000000 | 361.000000 | 361.000000 |
| mean | 0.383701 | 0.350081 | 0.390582 | 0.127424 | 0.268698 | 0.875346 | 0.168975 | 0.349000 |
| std | 0.209712 | 0.207184 | 0.149146 | 0.224465 | 0.287833 | 0.330784 | 0.375250 | 0.477000 |
| min | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 |
| 25% | 0.237143 | 0.189829 | 0.200000 | 0.000000 | 0.000000 | 1.000000 | 0.000000 | 0.000000 |
| 50% | 0.338095 | 0.295092 | 0.400000 | 0.000000 | 0.333333 | 1.000000 | 0.000000 | 0.000000 |
| 75% | 0.514286 | 0.491425 | 0.400000 | 0.000000 | 0.333333 | 1.000000 | 0.000000 | 1.000000 |
| max | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 1.000000 |



```
plt.figure(figsize = (16, 10))
sns.heatmap(df_train.corr(), annot = True, cmap="YlGnBu")
plt.show()
```

| | | | | | | | | | | | | | | |
|-----------------|-------|--------|---------|--------|--------|--------|--------|--------|---------|------|--------|--------|-------|--------|
| price | 1 | 0.56 | 0.3 | 0.42 | 0.43 | 0.28 | 0.25 | 0.2 | 0.059 | 0.47 | 0.32 | 0.33 | 0.13 | -0.28 |
| area | 0.56 | 1 | 0.14 | 0.19 | 0.11 | 0.25 | 0.25 | 0.11 | -0.025 | 0.26 | 0.33 | 0.21 | 0.053 | -0.12 |
| bedrooms | 0.3 | 0.14 | 1 | 0.32 | 0.4 | -0.013 | 0.029 | 0.1 | -0.0066 | 0.15 | 0.13 | 0.087 | 0.098 | -0.14 |
| bathrooms | 0.42 | 0.19 | 0.32 | 1 | 0.26 | 0.0088 | 0.17 | 0.12 | 0.046 | 0.2 | 0.15 | 0.029 | 0.039 | -0.12 |
| stories | 0.43 | 0.11 | 0.4 | 0.26 | 1 | 0.13 | -0.019 | -0.2 | 0.012 | 0.32 | 0.0036 | 0.052 | 0.022 | -0.033 |
| mainroad | 0.28 | 0.25 | -0.013 | 0.0088 | 0.13 | 1 | 0.081 | 0.048 | 0.032 | 0.11 | 0.17 | 0.17 | 0.058 | -0.11 |
| guestroom | 0.25 | 0.25 | 0.029 | 0.17 | -0.019 | 0.081 | 1 | 0.35 | 0.024 | 0.11 | 0.025 | 0.22 | 0.043 | -0.12 |
| basement | 0.2 | 0.11 | 0.1 | 0.12 | -0.2 | 0.048 | 0.35 | 1 | 0.0034 | 0.07 | 0.09 | 0.27 | 0.035 | -0.12 |
| hotwaterheating | 0.059 | -0.025 | -0.0066 | 0.046 | 0.012 | 0.032 | 0.024 | 0.0034 | 1 | -0.1 | 0.063 | -0.067 | 0.096 | -0.071 |
| airconditioning | 0.47 | 0.26 | 0.15 | 0.2 | 0.32 | 0.11 | 0.11 | 0.07 | -0.1 | 1 | 0.12 | 0.11 | -0.01 | -0.12 |

```
y_train = df_train.pop('price')
X_train = df_train
```



```
y_test = df_test.pop('price')
X_test = df_test
```

price area bedrooms bathrooms stories mainroad guestroom basement hotwaterheating airconditioning

```
X_train = X_train.to_numpy()
y_train = y_train.to_numpy()
```

```
from sklearn.neighbors import KNeighborsRegressor
```

```
knn = KNeighborsRegressor(n_neighbors=6, weights='uniform')
knn.fit(X_train,y_train)
```

```
KNeighborsRegressor(n_neighbors=6)
```

```
def apply_knn(neigh, weight='uniform'):
    knn = KNeighborsRegressor(n_neighbors=neigh, weights=weight)
    knn.fit(X_train,y_train)
    pred_knn = knn.predict(X_test)
    return pred_knn
```

```
model = KNeighborsRegressor()
params = {'n_neighbors':list(range(1, 50, 2)), 'weights':['uniform', 'distance']}
gs = GridSearchCV(model, params, cv = 5, n_jobs=-1)
gs_results = gs.fit(X_train, y_train)
```

```
print('Best Accuracy: ', gs_results.best_score_)
print('Best Paramtrs: ', gs_results.best_params_)
```

```
Best Accuracy: 0.523383914449
Best Paramtrs: {'n_neighbors': 9, 'weights': 'distance'}
```

```
pred_knn = apply_knn(9)
from sklearn.metrics import mean_absolute_error
print("MAE",mean_absolute_error(y_test,pred_knn))
```

```
MAE 4565790.77872
/usr/local/lib/python3.8/dist-packages/sklearn/base.py:443: UserWarning: X has feature names,
  warnings.warn(
```



```
from sklearn.metrics import mean_squared_error
print("RMSE",np.sqrt(mean_squared_error(y_test,pred_knn)))
```

```
RMSE 4876137.31684
```

```
print(knn.score(X_test,y_test))
```

```
-7.11417318899
/usr/local/lib/python3.8/dist-packages/sklearn/base.py:443: UserWarning: X has feature names,
  warnings.warn(
```



```
from sklearn import linear_model
regr = linear_model.LinearRegression()
```

```
regr.fit(X_train, y_train)
```

```
LinearRegression()
```


```
y_train_price = regr.predict(X_train)
```

```
res = (y_train_price - y_train)
```


```
fig = plt.figure()
sns.distplot((y_train - y_train_price), bins = 20)
fig.suptitle('Error Terms', fontsize = 20)
plt.xlabel('Errors', fontsize = 18)
```

```
/usr/local/lib/python3.8/dist-packages/seaborn/distributions.py:2619: FutureWarning: `distplot`
warnings.warn(msg, FutureWarning)
Text(0.5, 0, 'Errors')
```

Error Terms



```
plt.scatter(y_train,res)
plt.show()
```



Q7

```
train=pd.read_csv("/content/train.csv")
```

```
test=pd.read_csv("/content/test(1).csv")
```

```
train.isnull().sum()
```

```
PassengerId      0
Survived          0
Pclass           0
Name             0
Sex              0
Age             177
SibSp            0
Parch           0
Ticket           0
Fare            0
Cabin          687
Embarked         2
dtype: int64
```

```
test.isnull().sum()
```

```
PassengerId      0
Pclass           0
Name             0
```

```
Sex          0
Age          86
SibSp        0
Parch        0
Ticket       0
Fare         1
Cabin       327
Embarked     0
dtype: int64
```

```
impute_value = train['Age'].median()
```

```
test['Age'] = test['Age'].fillna(impute_value)
```

```
train['Age'] = train['Age'].fillna(impute_value)
```

```
train['IsFemale'] = (train['Sex'] == 'female').astype(int)
```

```
test['IsFemale'] = (test['Sex'] == 'female').astype(int)
```

```
predictors = ['Pclass', 'IsFemale', 'Age']
```

```
X_train = train[predictors].values
```

```
X_train
```

```
array([[ 3.,  0., 22.],
       [ 1.,  1., 38.],
       [ 3.,  1., 26.],
       ...,
       [ 3.,  1., 28.],
       [ 1.,  0., 26.],
       [ 3.,  0., 32.]])
```

```
X_test = test[predictors].values
```

```
X_test
```

```
array([[ 3. ,  0. , 34.5],
       [ 3. ,  1. , 47. ],
       [ 2. ,  0. , 62. ],
       ...,
       [ 3. ,  0. , 38.5],
       [ 3. ,  0. , 28. ],
       [ 3. ,  0. , 28. ]])
```

```
y_train = train['Survived'].values
```

```
y_train
```

```
array([0, 1, 1, 1, 0, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 1,
       1, 0, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1, 1, 0,
       0, 1, 0, 0, 0, 0, 1, 1, 0, 1, 1, 0, 1, 0, 0, 1, 0, 0, 0, 1, 1, 0, 1,
       0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 1, 1, 0, 1, 1, 0, 0, 1, 0, 0, 0,
       0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0,
       0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0,
```

```

0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0,
1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1,
1, 0, 1, 1, 0, 0, 1, 0, 1, 1, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0,
1, 1, 1, 0, 1, 0, 0, 0, 1, 1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0,
1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0,
0, 0, 1, 1, 1, 1, 1, 0, 1, 0, 0, 0, 0, 0, 1, 1, 1, 0, 1, 1, 0, 1, 1,
0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1,
1, 1, 1, 0, 1, 0, 1, 1, 1, 0, 1, 1, 1, 0, 0, 0, 1, 1, 0, 1, 1, 0, 0,
1, 1, 0, 1, 0, 1, 1, 1, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 0,
1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 1,
1, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 1, 0, 1, 1,
1, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0,
1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 1, 1, 1, 0, 0, 1, 0,
1, 0, 0, 1, 0, 0, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 1, 0, 1, 0, 1, 1, 0,
1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0,
1, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0,
1, 1, 0, 1, 1, 0, 1, 1, 0, 0, 1, 0, 1, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0,
0, 1, 0, 0, 1, 0, 1, 0, 1, 0, 1, 1, 0, 0, 1, 0, 0, 1, 1, 0, 1, 1, 0,
0, 1, 1, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 0,
0, 1, 1, 0, 1, 1, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0,
0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0,
1, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 1,
1, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0,
0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1,
1, 1, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 0, 1,
0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 1, 1,
0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 1,
1, 1, 1, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0,
0, 0, 1, 0, 1, 1, 1, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1, 0, 1, 0, 0,
1, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0])

```

```

X_train[:5]

```

```

array([[ 3.,  0., 22.],
       [ 1.,  1., 38.],
       [ 3.,  1., 26.],
       [ 1.,  1., 35.],
       [ 3.,  0., 35.]])

```

```

y_train[:5]

```

```

array([0, 1, 1, 1, 0])

```

```

from sklearn.linear_model import LogisticRegression
model = LogisticRegression()
model.fit(X_train, y_train)

```

```

LogisticRegression()

```

```

y_predict = model.predict(X_test)
y_predict[:10]

```

```

array([0, 0, 0, 0, 1, 0, 1, 0, 1, 0])

```

test

| | PassengerId | Pclass | Name | Sex | Age | SibSp | Parch | Ticket | Fare | Cabin | Embarked |
|-----|-------------|--------|--|--------|------|-------|-------|-----------|---------|-------|----------|
| 0 | 892 | 3 | Kelly, Mr. James | male | 34.5 | 0 | 0 | 330911 | 7.8292 | NaN | |
| 1 | 893 | 3 | Wilkes, Mrs. James (Ellen Needs) | female | 47.0 | 1 | 0 | 363272 | 7.0000 | NaN | |
| 2 | 894 | 2 | Myles, Mr. Thomas Francis | male | 62.0 | 0 | 0 | 240276 | 9.6875 | NaN | |
| 3 | 895 | 3 | Wirz, Mr. Albert | male | 27.0 | 0 | 0 | 315154 | 8.6625 | NaN | |
| 4 | 896 | 3 | Hirvonen, Mrs. Alexander (Helga E Lindqvist) | female | 22.0 | 1 | 1 | 3101298 | 12.2875 | NaN | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 413 | 1305 | 3 | Spector, Mr. Woolf | male | 28.0 | 0 | 0 | A.5. 3236 | 8.0500 | NaN | |

```
from sklearn.linear_model import LogisticRegressionCV
model_cv = LogisticRegressionCV()
model_cv.fit(X_train, y_train)

LogisticRegressionCV()

from sklearn.model_selection import cross_val_score
model = LogisticRegression(C=10)
scores = cross_val_score(model, X_train, y_train, cv=4)
scores

array([ 0.77578475,  0.79820628,  0.77578475,  0.78828829])
```